

Industry can get any empirical
research it wants

(Publish open source data, and some example scripts.)

Tim Menzies

prof, cs, **ncstate**, usa
acm-ieee-ase fellow; eic ASEj
timm@ieee.org
<http://timm.fyi>

Oct3'25



[Meno07]: Data Mining Static Code Attributes to Learn Defect Predictors,
TSE'07

[Men25] T. Menzies, "Retrospective: Data Mining Static Code Attributes, TSE'25

The Portland Context

- Born from open source culture in Portland, Oregon
- *"We wore no suite and tie in our photos. We did not comb our hair"*
- Philosophy: `svn commit -m "share stuff"` will change SE research
- But unhappy with SOTA data mining in SE
- **Key Insight:** Walking around Chicago's Grant Park (2004)
 - **Tim Menzies** and **Jelber Sayyad** lamented: *"Must do better... Why don't we make conclusions reproducible?"*

The Radical Idea

- In 2025 hard to believe “reproducible SE” was radical
- **Lionel Briand** (2006): *“no one will give you data”*
- Yet we persisted...

2005? Really? What have you done since?

What do I do? How can I help you?

papers, 9 | are-articles, 2 | data, 4 | drexelis.org | 1-334-376-2839 | 111 334, E20 [exp]

Tim Menzies

<http://timm.fyi>

"AI, but clearer."

So I seek talented grad students & industrial partners to find + fix the problems in real-world AI&ML.
Is that you? Maybe "yes" if you want to be a leader in AI (and not just another follower).

[Job Advertisements](#)

videos

total funding
(in 2025 dollars, assuming 3% inflation)

Year	Funding (Millions)
1995	0.5
2000	1.0
2005	2.5
2010	5.0
2015	8.0
2020	12.0
2025	18.0

Career:

- Bookkeeping: Graduate Miles to 30
- Editor-in-chief: Automated Software Engineering journal
- Associate Editor: IEEE Trans SE
- Editorial Board: Communications of the ACM (opinions)
- Third author: 100+ papers, 100+ programs
- Program committees: ICSE'26, ICSE'25, FSE'25, AIRE'25, SANER'25, ESEM'25, CAIN'25, AAAI'25, ICSE'23

Ph.D. (current): Ph.D. (graduation):

Much thanks and good luck to you!



National
Science
Foundation



Last decade: ~140 papers



Major Commercial Breakthrough Technologies:

- **NASA-Perkins Analytics Platform** - Mission-critical systems, space-grade reliability (27 papers)
 - **Microsoft Hyperparameter Optimization** - 210X faster, enterprise validated (12 papers)
 - **Microsoft Text Mining Platform** - Legal document analysis, millions of documents
 - **Cloud Configuration Systems** - Prem 2 cloud (4) / cloud config (4) solutions
 - **Software Design Intelligence** - AI-driven design optimization (7 papers)
 - **Project Health Analytics** - Predictive monitoring and early warning systems (6 papers)

Two-Part Vision:

- ➊ **Annual conference** on predictor models in SE (to share results)
- ➋ **Repository** of 100s of SE datasets: defect prediction, effort estimation, Github issue close time, bad smell detection

Growth Trajectory:

- Repository grew; moved to **Large Hadron Collider** (Seacraft, Zenodo)
- Research students ran weekly sprints scouring SE conferences
- Gary Boetticher, Elaine Weyuker, Thomas Ostrand, Guenther Ruhe** joined steering committee → prestige for growth

PROMISE vs MSR:

- MSR**: Gathering initial datasets (**Devanbu [Dev15]**)
- PROMISE**: Post-collection analysis, data re-examination **[Rob10]**

Early Results:

- Other areas struggled with reproducibility, while we swam in data
- Papers applied tool sets to COC81, JM1, XALAN, DESHARNIS etc
- First decade: Numerous successful papers using consistent data re-examination

Research Question: Can data mining algorithms learn software defect predictors from static code attributes?

Why This Matters:

- “Software quality assurance budgets are finite while assessment effectiveness increases exponentially with effort” **[Fu16]**
- “Software bugs are not evenly distributed across a project” **[Ham09], [Osto4], [Mis11]**
- Defect predictors suggest where to focus expensive methods

Counter-Arguments Addressed:

- ① “Specific metrics matter” (1990s heated debates: McCabe vs Halstead)
- ② “Static code attributes do not matter” (**Fenton & Pfleeger, Shepperd & Ince**)

1st Law: "Specific metrics do not always matter in all data sets. Rather, different projects have different best metrics."

Supporting Evidence:

- Feature pruning experiment on **3 dozen metrics across 7 datasets**
- Results: Pruning selected just **2-3 attributes per dataset**
- **No single attribute** selected by majority of datasets
- Different projects preferred different metrics (McCabe vs Halstead vs lines of code)
- Theoretical debates of 1990s (metric X vs metric Y) proven empirically unfounded

Menzies's Corollary:

"To mine SE data, gather all that can be collected (cheaply) then apply data pruning to discard irrelevancies."

Practical Impact:

- Changed SE data mining methodology from “careful metric selection” to “gather everything, prune later”

2nd Law: “Static code attributes do matter. Individually, they may be weak indicators. But when combined, they can lead to strong signals that outperform the state-of-the-art.”

Support Evidence:

- **Fenton & Pfleeger:** Same functionality, different constructs → different measurements
- **Shepperd & Ince:** Static measures often “no more than proxy for lines of code”
- **Our Response:** Stress-tested these views by documenting baselines, then showing detectors from static attributes **much better** than baselines
- **Key Finding:** Multi-attribute models outperformed single-attribute models

Key Quote: “Paradoxically, this paper will be a success if it is quickly superseded.”

Citation Impact:

- **2016**: Most cited paper (per month) in software engineering
- **2018**: 20% of Google Scholar Software Metrics IEEE TSE papers used PROMISE datasets [**Meno7**]
- **Current**: 1924 citations (paper) + 1242 citations (repository)

Industrial Adoption:

- **Wan et al. [Wan20]**: 90%+ of 395 commercial practitioners willing to adopt defect prediction
- **Misirli et al. [Mis11]**: 87% defect prediction accuracy, 72% reduced inspection effort, 44% fewer post-release defects
- **Kim et al. [Kim15]**: Samsung Electronics API development
 - 0.68 F1 scores, reduced test case resources

Rahman et al. [Rah14] Comparison:

- **Static analysis tools:** FindBugs, Jlint, PMD
- **Statistical defect prediction:** Logistic regression models
- **Result:** “No significant differences in cost-effectiveness were observed”

Critical Advantage:

- Defect prediction: Quick adaptation to new languages via lightweight parsers
- Static analyzers: Extensive modification required for new languages
- **Implication:** Broader applicability across programming ecosystems

Extended Applications:

- **Security vulnerabilities** [Shi13]
- **Resource allocation** for defect location [Bir21]
- **Proactive defect fixing** [Kam16], [LeG12], [Arc11]
- **Change-level/just-in-time prediction** [Yan19], [Kam13], [Nay18], [Ros15]
- **Transfer learning** across projects [Kri19], [Nam18]
- **Hyperparameter optimization** [Agr18], [Che18], [Fu17], [Tan16]

Research Evolution:

- From binary classification to multi-objective optimization
- From release-level to line-level prediction (**Pornprasit et al. [Por23]** - TSE Best Paper 2023)

Phase Evolution:

- ① “Data? Good luck with that!” - Resistance and skepticism
- ② “Okay, maybe it’s not completely useless.” - Grudging acknowledgment
- ③ “This is the gold standard now.” - Required baseline, field norms
- ④ “A graveyard of progress.” - Stifling creativity, outdated paradigms

The Problem:

- Decade 2: Continued use of decades old data e.g. COC81 (1981), DESHARNIS (1988), JM1 (2004), XALAN (2010)
- **Editorial Policy Change:** Automated Software Engineering journal now desk-rejects papers based on 2005 datasets

3rd law: “Turkish toasters can predict errors in deep space satellites.”

Supporting Evidence:

- **Transfer learning research [Turo9]:** Models from **Turkish white goods** successfully predicted errors in **NASA systems**
- Expected: Complex multi-dimensional transforms mapping attributes across domains
- **Reality:** Simple nearest neighboring between test and training data worked perfectly
- **Implication:** “*Many distinctions made about software are spurious and need to be revisited*”

Broader Transfer Learning Success:

- Cross-domain prediction often works better than expected
- Suggests universal patterns in software defect manifestation
- Questions assumptions about domain-specific modeling requirements

4th Law: "For SE, the best thing to do with most data is to throw it away."

Supporting Evidence:

- **Chen, Kocaguneli, Tu, Peters, and Xu et al.** findings across multiple prediction tasks:
 - **Github issue close time:** Ignored 80% of data labels [**Che19**]
 - **Effort estimation:** Ignored 91% of data [**Koc13**]
 - **Defect prediction:** Ignored 97% of data [**Pet15**]
 - **Some tasks:** Ignored 98-100% of data [**Che05**]
- **Startling result:** Data sets with thousands of rows modeled with just few dozen samples [**Meno08**]

Theoretical Explanations:

- **Power laws** in software data [**Lin15**]
- **Large repeated structures** in SE projects [**Hin12**]
- **Manifold assumption** and **Johnson-Lindenstrauss lemma** [**Zhuo05**], [**Joh84**]

Caveat: Applies to regression, classification, optimization

- generative tasks may still need massive data

5th law: “Bigger is not necessarily better.”

Supporting Evidence - LLM Hype Analysis:

- **Systematic review [Hou24]:** 229 SE papers using Large Language Models
- **Critical finding:** Only **13/229 around 5%** compared LLMs to other approaches
- “*Methodological error*” - other PROMISE-style methods often better/faster **[Gri22], [Som24], [Taw23], [Maj18]**

Trading Off Complexity:

- Scalability vs. privacy vs. performance **[Lin24], [Fu17]**
- Often simpler methods provide better cost-effectiveness
- **Personal Pattern:** “Often, I switch to the simpler.” **[Agr21], [Tan16], [Fu16]**

6th Law: “Data quality matters less than you think.”

Supporting Research:

- **Shepperd et al. [She13]:** Found numerous PROMISE data quality issues
 - Repeated rows, illegal attributes, inconsistent formats
 - **Critical gap:** Never tested if quality issues decreased predictive power

Our Experiment:

- Built **mutators** that injected increasing amounts of their quality issues into PROMISE defect datasets
- **Startling result:** Performance curves remained **flat** despite increased quality problems
- **Implication:** “*There is such a thing as too much care*” in data collection

Practical Impact:

- Effective predictions possible from seemingly dirty data
- Questions excessive data cleaning efforts in SE research
- Balance needed: careful collection without over-engineering

7th Law: "Bad learners can make good conclusions."

Supporting Evidence:

- **Nair et al. [Nai17]**: CART trees built for multi-objective optimization
- **Key finding**: Models that **predicted poorly** could still **rank solutions effectively**
- Could be used to prune poor configurations and find better ones
- **Implication**: Algorithms shouldn't aim for predictions but offer **weak hints** about project data

8th Law: "Science has mud on the lens."

Supporting Evidence:

- **Hyperparameter optimization** lessons [**Agr21**], [**Tan16**], [**Fu16**] on PROMISE data
- Data mining conclusions **changeable in an afternoon** by grad student with sufficient CPU
- **Critical Questions:** Are all conclusions brittle? How build scientific community on such basis?
- **Where are stable conclusions** for building tomorrow's ideas?

?Bayesian Approach Needed: Address uncertainty quantification and robust foundations

9th Law: "Many hard SE problems, aren't."

Supporting Philosophy:

- **Cohen's Straw Man Principle [Coh95]:** "Supposedly sophisticated methods should be benchmarked against seemingly stupider ones"

Personal Experience Pattern:

- "Whenever I checked a supposedly sophisticated method against a simpler one, there was always something useful in the simpler"
- "Often, I switch to the simpler." **[Agr21], [Tan16], [Fu16]**

Important Caveat:

- **Not all SE problems can/should be simplified** (safety-critical; generative);
- "Just because some tasks are hard, does not mean all tasks are hard"

Challenge to Community: "Have we really checked what is really complex and what is really very simple?"

Current Focus: Minimal data approaches - landscape analysis **[Che19], [Lus24]**, surrogate learning **[Nai20]**, active learning **[Kra15], [Yu18]**

PROMISE Revival Strategy (**Gema Rodríguez-Pérez**):

- Data sharing now expected for almost all SE papers
- PROMISE must differentiate: accept higher quality datasets
- Focus on enhancing current data space, conducting quality evaluations

Steffen Herbold's Caution:

- Early PROMISE: Collections of metrics (not raw data)
- MSR shift: Raw data + fast tools (e.g., PyDriller, GHTorrent)
- **Risk:** “*Little curation, little validation, often purely heuristic data collection without quality checks*” [**Her22**]

Modern Data Access: 1100+ recent Github projects [**Xia22**], CommitGuru [**Ros15**]

Contemporary Approaches:

- **DeepLineDP (Pornprasit et al. [Por23])**: Deep learning for line-level defect prediction (TSE Best Paper 2023)
- **Model interpretability**: Growing research focus [**Tan21**]
- **Multi-objective optimization**: Hyperparameter selection [**Xia22**], unfairness reduction [**Cha20**], [**Alv23**]

Optimize CPU-Intensive Algorithms:

- MaxWalkSat [**Meno9**]
- Simulated annealing [**Meno2**], [**Meno7**]
- Genetic algorithms

Minimal Data Approaches:

- How much can be achieved with as little data as possible?
- Suspicion of “large number of good quality labels” assumption

Cross-Domain Success [Turo9]:

- **Turkish white goods → NASA systems** error prediction
- Expected: Complex multi-dimensional transforms
- **Reality:** Simple nearest neighboring between test and training data

Implication: “Many distinctions made about software are spurious and need to be revisited”

Power Laws & Repeated Structures:

- **Lin & Whitehead [Lin15]:** Fine-grained code changes follow power laws
- **Hindle et al. [Hin12]:** Software naturalness - large repeated structures
- **Result:** Thousands of rows modeled with few dozen samples **[Meno8]**

Lessons Learned:

- ① **Open science communities** can be formed by publishing baseline + data + scripts
- ② **Reproducible research** drives field advancement when embraced collectively
- ③ **Simple solutions** often outperform sophisticated ones
- ④ **Data quality** matters less than expected for predictive tasks
- ⑤ **Transfer learning** works across surprisingly diverse domains

Call-to-Action:

- “Have we really checked what is really complex and what is really very simple?”
- Challenge assumptions about problem complexity
- Benchmark sophisticated methods against simpler alternatives
- Focus on stable, reproducible conclusions

- [Agr18]:** A. Agrawal and T. Menzies, “Is better data better than better data miners?: On the benefits of tuning smote for defect prediction,” in *Proc. IST*, ACM, 2018, pp. 1050–1061.
- [Agr21]:** A. Agrawal *et al.*, “How to”DODGE” complex software analytics?” *IEEE Trans. Softw. Eng.*, vol. 47, no. 10, pp. 2182–2194, Oct. 2021.
- [Alv23]:** L. Alvarez and T. Menzies, “Don’t lie to me: Avoiding malicious explanations with STEALTH,” *IEEE Softw.*, vol. 40, no. 3, pp. 43–53, May/Jun. 2023.
- [Cha20]:** J. Chakraborty *et al.*, “Fairway: A way to build fair ML software,” in *Proc. FSE*, 2020, pp. 654–665.
- [Che19]:** J. Chen *et al.*, “‘Sampling’ as a baseline optimizer for search-based software engineering,” *IEEE Trans. Softw. Eng.*, vol. 45, no. 6, pp. 597–614, Jun. 2019.
- [Coh95]:** P. R. Cohen, *Empirical Methods for Artificial Intelligence*, Cambridge, MA: MIT Press, 1995.
- [Dev15]:** P. Devanbu, “Foreword,” in *Sharing Data and Models in Software Engineering*, T. Menzies *et al.*, Eds. San Mateo, CA: Morgan Kaufmann, 2015, pp. vii–viii.
- [Fu16]:** W. Fu, T. Menzies, and X. Shen, “Tuning for software analytics: Is it really necessary?” *Inform. Softw. Technol.*, vol. 76, pp. 135–146, 2016.

- [Gon23]:** J. M. Gonzalez-Barahona and G. Robles, "Revisiting the reproducibility of empirical software engineering studies based on data retrieved from development repositories," *Inf. Softw. Technol.*, vol. 164, 2023, Art. no. 107318.
- [Gri22]:** L. Grinsztajn, E. Oyallon, and G. Varoquaux, "Why do tree-based models still outperform deep learning on typical tabular data?" in *Proc. NeurIPS*, 2022, pp. 507–520.
- [Ham09]:** M. Hamill and K. Goseva-Popstojanova, "Common trends in software fault and failure data," *IEEE Trans. Softw. Eng.*, vol. 35, no. 4, pp. 484–496, Jul./Aug. 2009.
- [Has08]:** A. E. Hassan, "The road ahead for mining software repositories," *Frontiers Softw. Maintenance*, pp. 48–57, 2008.
- [Her22]:** S. Herbold *et al.*, "Problems with SZZ and features: An empirical study of the state of practice of defect prediction data collection," *Empirical Softw. Eng.*, vol. 27, no. 2, p. 42, 2022.
- [Hou24]:** X. Hou *et al.*, "Large language models for software engineering: A systematic literature review," *ACM Trans. Softw. Eng. Methodol.*, vol. 33, no. 8, Dec. 2024.
- [Kam13]:** Y. Kamei *et al.*, "A large-scale empirical study of just-in-time quality assurance," *IEEE Trans. Softw. Eng.*, vol. 39, no. 6, pp. 757–773, Jun. 2013.
- [Kim15]:** M. Kim *et al.*, "REMI: Defect prediction for efficient api testing," in *Proc. FSE*, ACM, 2015, pp. 990–993.

References (Yet More)

- [Kri19]:** R. Krishna and T. Menzies, "Bellwethers: A baseline method for transfer learning," *IEEE Trans. Softw. Eng.*, vol. 45, no. 11, pp. 1081–1105, Nov. 2019.
- [Meno07]:** T. Menzies, J. Greenwald, and A. Frank, "Data mining static code attributes to learn defect predictors," *IEEE Trans. Softw. Eng.*, vol. 33, no. 1, pp. 2–13, Jan. 2007.
- [Men24]:** T. Menzies, "A brief note, with thanks, on the contributions of guenther ruhe," *Inf. Softw. Technol.*, vol. 173, 2024, Art. no. 107486.
- [Men25]:** T. Menzies, "Retrospective: Data Mining Static Code Attributes to Learn Defect Predictors," *IEEE Trans. Softw. Eng.*, 2025.
- [Mis11]:** A. T. Misirli, A. Bener, and R. Kale, "AI-based software defect predictors: Applications and benefits in a case study," *AI Mag.*, vol. 32, no. 2, pp. 57–68, 2011.
- [Nai17]:** V. Nair *et al.*, "Using bad learners to find good configurations," in *Proc. 11th Joint Meeting FSE*, ACM, 2017, pp. 257–267.
- [Nam18]:** J. Nam *et al.*, "Heterogeneous defect prediction," *IEEE Trans. Softw. Eng.*, vol. 44, no. 9, pp. 874–896, Sep. 2018.
- [Osto04]:** T. J. Ostrand, E. J. Weyuker, and R. M. Bell, "Where the bugs are," *ACM SIGSOFT Softw. Eng. Notes*, vol. 29, no. 4, pp. 86–96, 2004.
- [Por23]:** C. Pornprasit and C. K. Tantithamthavorn, "DeepLineDP: Towards a deep learning approach for line-level defect prediction," *IEEE Trans. Softw. Eng.*, vol. 49, no. 1, pp. 84–98, Jan. 2023.
- [Rah14]:** F. Rahman *et al.*, "Comparing static bug finders and statistical prediction," in *Proc. ICSE*, ACM, 2014, pp. 424–434.

- [Rob10]:** G. Robles, "Replicating MSR: A study of the potential replicability of papers published in the mining software repositories proceedings," in *7th IEEE Work. Conf. Mining Softw. Repositories (MSR)*, IEEE Press, 2010, pp. 171–180.
- [Ros15]:** C. Rosen, B. Grawi, and E. Shihab, "Commit guru: Analytics and risk prediction of software commits," in *Proc. ESEC/FSE*, 2015, pp. 966–969.
- [She13]:** M. Shepperd *et al.*, "Data quality: Some comments on the NASA software defect datasets," *IEEE Trans. Softw. Eng.*, vol. 39, no. 9, pp. 1208–1215, Sep. 2013.
- [Shi13]:** Y. Shin and L. Williams, "Can traditional fault prediction models be used for vulnerability prediction?" *Empirical Softw. Eng.*, vol. 18, pp. 25–59, 2013.
- [Tan16]:** C. Tantithamthavorn *et al.*, "Automated parameter optimization of classification techniques for defect prediction," in *ICSE'16*, 2016, pp. 321–332.
- [Tur09]:** B. Turhan *et al.*, "On the relative value of cross-company and within-company data for defect prediction," *Empirical Softw. Eng.*, vol. 14, pp. 540–578, Jan. 2009.
- [Wan20]:** Z. Wan *et al.*, "Perceptions, expectations, & challenges in defect prediction," *IEEE Trans. Softw. Eng.*, vol. 46, no. 11, pp. 1241–1266, Nov. 2020.
- [Xia22]:** T. Xia *et al.*, "Sequential model optimization for software effort estimation," *IEEE Trans. Softw. Eng.*, vol. 48, no. 6, pp. 1994–2009, Jun. 2022.
- [Yan19]:** M. Yan *et al.*, "Automating change-level self-admitted technical debt determination," *IEEE Trans. Softw. Eng.*, vol. 45, no. 12, pp. 1211–1229, Dec. 2019.